

Deep Deterministic Policy Gradient

—

Anirudh Goyal, Rosemary Nan Ke

Continuous actions ?

- Most interesting problems fall in this category.
- If you discretize your action space naively,
 - Curse of dimensionality problem as before.
 - Throws away valuable information regarding the geometry of action domain

Policy gradient

- Instead of estimating the value or action-value function.
Directly parameterize the policy
- Used for continuous action spaces
- Stochastic policy $\mu(a|s)$ - probability over actions
- Seems many examples with high reward for good actions,
and many examples with negative reward for bad actions.

Policy gradient


- Optimize the policy end-to-end by computing noisy estimates of the gradient of the expected reward of the policy and then updating the policy in the gradient direction.
- Traditionally, PG methods have assumed a stochastic policy.

Policy gradient

- Issues:
 - Getting one reward signal at the end of the episode, of interaction makes it impossible to ascertain which action was good (Credit assignment problem)
 - Curse of dimensionality

Actor Critic Algorithms

- Policy function - Actor , Value function - critic
- Actor: Input state s , output action a
- Critic: Input state s , reward r , output TD error difference



Off-Policy v/s On-Policy

Off-Policy

- Generally employ a separate behaviour policy that is independent of policy being improved upon, behaviour policy is used to simulate trajectories.
- **Key benefit - Behaviour policy can operate by sampling all actions, whereas estimation policy can be greedy (i.e deterministic)**
- Q policy is an off policy algorithm
- On policy algorithms directly use the policy that is being estimated to sample trajectories during training.

Model free algorithms

- No effort to learn the underlying dynamics that govern how an agent interacts with environment.
- In case, environment \rightarrow discrete state space, agent \rightarrow discrete number of actions, a model of the dynamics of the environment is 1 step transition operator. $T(s(t+1) \mid s(t), a(t))$
- For high dimensional state space, this is extremely expensive!!
- Hence, Model free algorithm directly estimate the policy or value function, (more computationally efficient)
- Be wary, Using a bad approximation of environment would bring you misery.

Deep Deterministic Policy Gradient

- Uses a stochastic behaviour policy for good exploration.
- Estimates a deterministic target policy.
- Utilize a form of policy iteration: they evaluate the policy and then follow policy gradient.
- Actor critic algorithm as well.
- In continuous space, finding the greedy policy requires an optimization of a_t at every time step. Use Actor critic based on DPG algorithm

Theorem. (Deterministic Policy Gradient Theorem). Suppose the Markov Decision Process satisfies the appropriate conditions (see [3]). These imply that $\nabla_{\theta^\mu} \mu(s|\theta^\mu)$ and $\nabla_a Q(s, a|\theta^Q)$ exist and that the deterministic policy gradient exists. Then,

$$\begin{aligned}\nabla_{\theta^\mu} \mu &\approx \int_S \rho^{\mu'}(s_t) \nabla_a Q(s, a|\theta^Q)|_{s=s_t, a=\mu(s_t)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s=s_t} ds \\ &= \mathbb{E}_{\mu'} [\nabla_a Q(s, a|\theta^Q)|_{s=s_t, a=\mu(s_t)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s=s_t}]\end{aligned}$$

Proof. For a greedy stochastic policy $\mu(a|s, \theta)$ over a continuous action space, a global maximization step is required at every time step. Rather, we employ a deterministic policy $\mu(s|\theta)$ and update the policy parameters by moving them in the direction of the gradient of the action-value function. We take an expectation to average over the suggested directions of improvement from each state w.r.t. the state distribution under the target policy μ' , given by $\rho^{\mu'}(s)$.

$$\theta_{k+1}^\mu = \theta_k^\mu + \alpha \mathbb{E}_{\mu'^k} [\nabla_{\theta} Q(s, \mu(s|\theta_k^\mu)|\theta_k^Q)].$$

By applying the chain rule,

$$\theta_{k+1}^\mu = \theta_k^\mu + \alpha \mathbb{E}_{\mu'^k} [\nabla_a Q(s, a|\theta_k^Q)|_{a=\mu(s|\theta_k^\mu)} \nabla_{\theta} \mu(s|\theta_k^\mu)].$$

DDPG

- Uses 2 networks
 - Actor - Input is current state, output is single real value representing the action chosen from a continuous action space.
 - Critic's output is simply the estimated Q value of the state and action given by actor (policy)
 - Deterministic policy gradient theory provides the update rule for the weights of the actor network.
 - Critic network is updated from the gradients obtained from the TD error signal.

Key Conspirators for divergence

- **Problem 1** - Training your policy for 1000's of temporally correlated simulated trajectories leads to the introduction of enormous amounts of variance in your approximation of true Q function. TD error signal is excellent in compounding the variance by bad predictions over time.
- **Problem - 2** - Directly updating the actor and critic weights with gradients obtained from TD error, can cause your learning algorithm to diverge.

Solutions - Lessons from Success of DQN

- NN for RL Assume, that samples are iid, but where the samples are generated from exploring sequentially in an environment, this assumption not holds.
- AS DQN, we use replay buffers, a technique called experience replay to address this issue.
- Using the set of target networks to generate the targets for your TD error computation regularizes your learning algorithm and increases stability.
- Targets network slowly change that greatly improve the stability of the training.

Challenge- 2

- Advantage of off policy algorithm is that you can treat the exploration problem independent of learning algorithm
- Construct an exploration policy by adding noise sampled from a noisy process.
- Use an ornstein-uhlenbeck process to generate temporally correlated exploration for exploration efficiency with inertia.

References

- Continuous control with deep reinforcement learning
- Deterministic Policy Gradient Algorithms
- Prioritized experience replay <https://arxiv.org/abs/1511.05952>